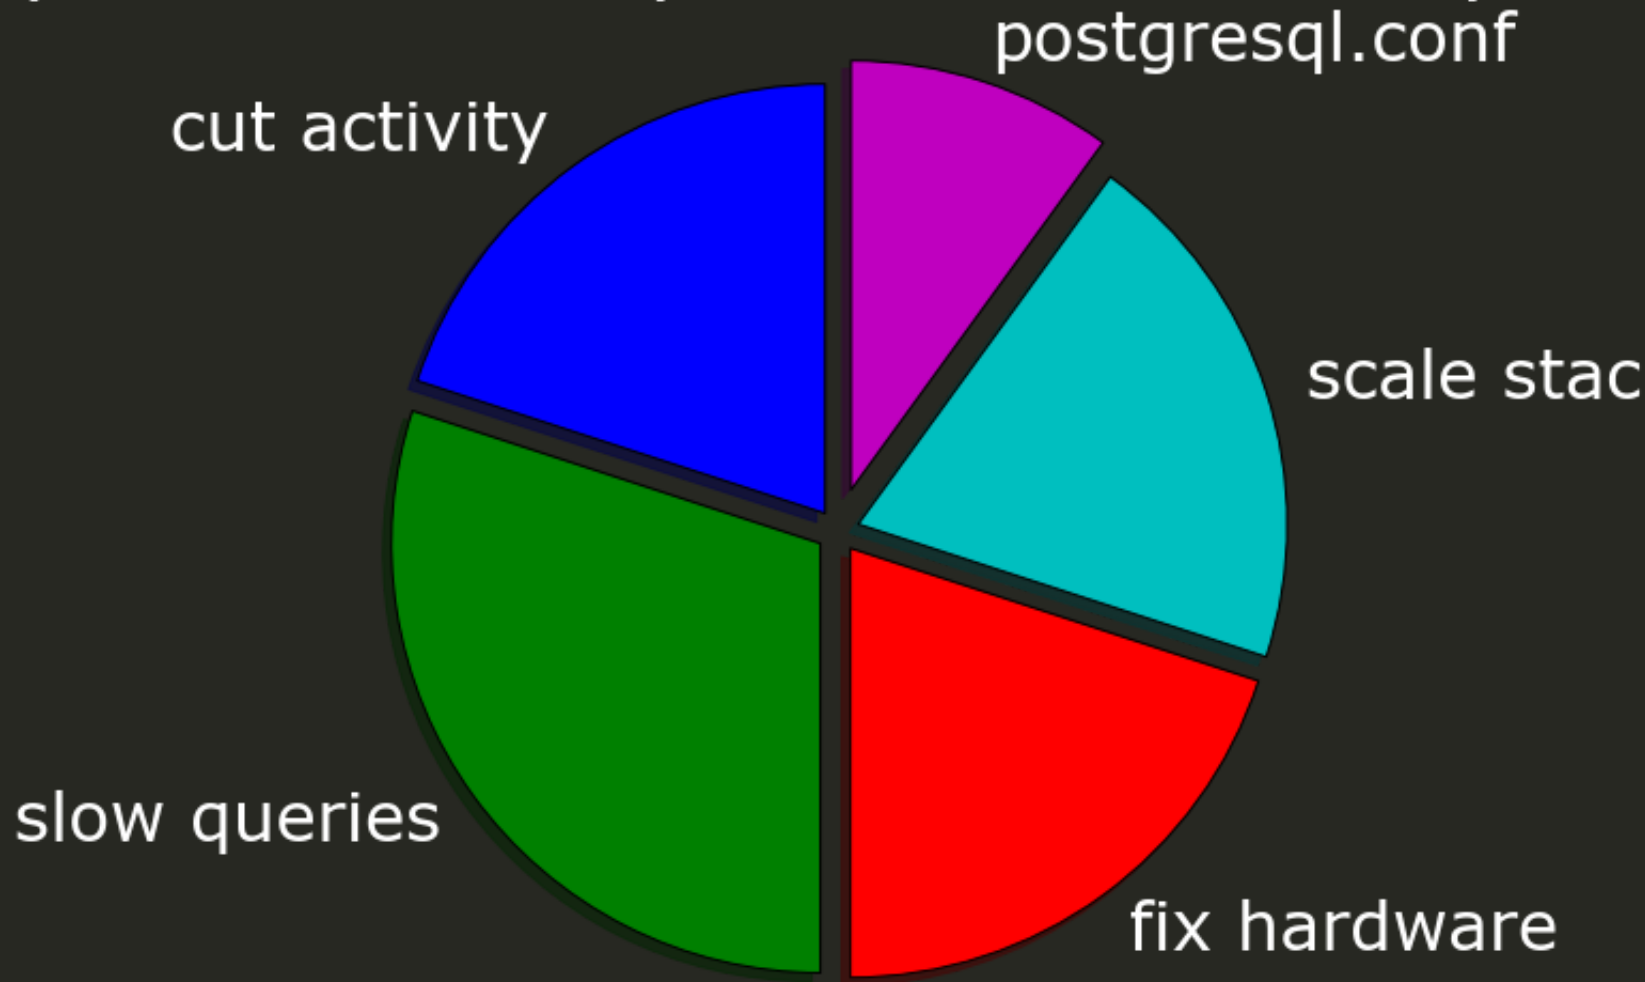


Postgres Performance in 15 Minutes




go_faster = 10

performance optimization: time spent



performance optimization: effect postgresql.conf



Do Less

do I need the
database to answer
that?

caching

- results cache
- redis
- memcached
- CDN

some data access
anti-patterns

polling

```
while True:  
    j = check_for_job(pid)  
    if j is None:  
        continue
```

data you already have

```
curUser = Users.object.filter(pk=sessionUser).values()  
curUserID = curUser[0]["id"]
```

data you already have

```
SELECT id FROM users  
WHERE id = ?
```

data you don't need

```
AllProfiles = Profile.objects.all().order_by('-updated')  
LastProfile = AllProfiles[0]
```

join loops

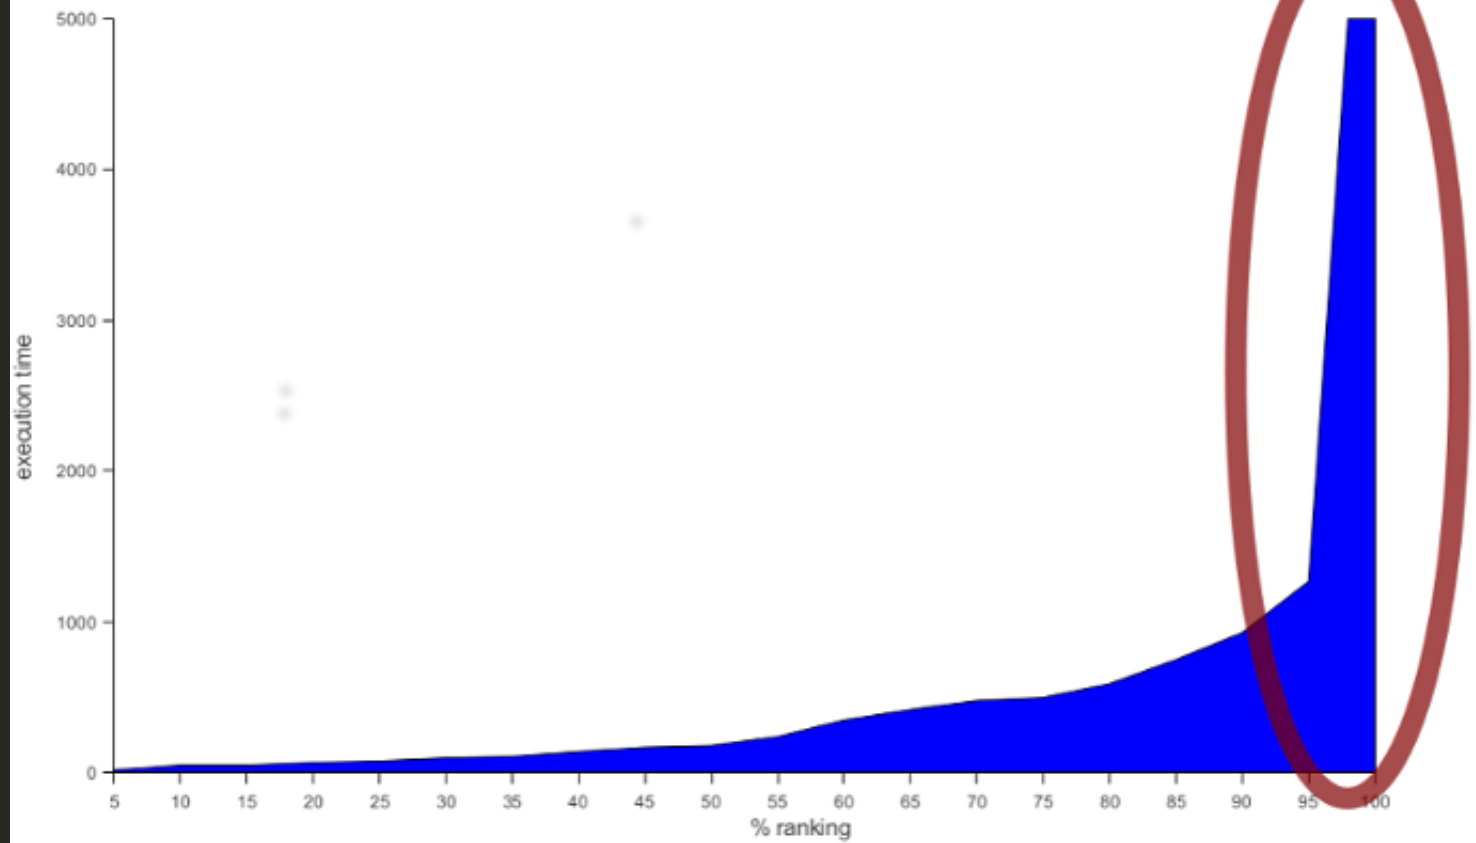
```
for Player in Roster:  
    myGames =  
        Games.object.filter(player_id=Player["id"]).values()
```

```
SELECT * FROM games WHERE player_id = ?  
-- repeat 200X
```



resource-hungry queries

Ranked Query Execution Times



SQL Traffic

KEY VALUES


406 queries/s

Query Peak

2012-12-08 21:40:04

Date

QUERIES PER SECOND (5 MINUTES AVERAGE)



To Image

Download

To Chart

pgBadger

⌚ Time consuming queries

Rank	Total duration	Times executed	Min duration	Max duration	Avg duration	Query
1	43s330ms	74 Details	304ms	1s690ms	585ms	🔗 <code>SELECT count (*) FROM pg_catalog.pg_class AS c WHERE c.oid = pgpool_regclass (' ') AND c.relpersistence = '';</code> Examples
2	40s841ms	1 Details	40s841ms	40s841ms	40s841ms	🔗 <code>DELETE FROM filterout WHERE last_update = ''; INSERT INTO filterout (script, os, arch, build, gpu, branch, testname, bug, active, category, vc, farm, osversion) SELECT script, os, arch, build, gpu, branch, testname, bug, active, category, vc, farm, osversion FROM filterout WHERE last_update = (SELECT max (last_update) FROM filterout WHERE last_update != now () ::date);</code> Examples
3	8s10ms	3 Details	446ms	4s543ms	2s670ms	🔗 <code>SELECT testname, category, bug FROM filterout WHERE os = '' AND osversion = '' AND arch = '' AND gpu = '' AND branch = '' AND build = '' AND vc = '' AND script = '' AND active = '' AND last_update = '' AND farm = '' ORDER BY testname;</code> Examples

pgBadger

improving slow queries

- add indexes
- fix filter expressions
- analyze/autoanalyze

text searching

startswith: requires varchar_pattern_ops

istartswith: requires function index or
Ctext

icontains: use TSearch or trigrams

transaction pitfalls

bad: bundling reads into a transaction

worse: holding a transaction open while rendering

worst: holding a transaction open while waiting for user input

locking

check for "waiting" queries:

```
SELECT * FROM pg_stat_activity WHERE waiting;
```

does your program logic have multiple threads update the same data at the same time?

adequate hardware

and virtual hardware



You can't outperform
inadequate hardware

IO is more important than
you think

postgres writes all the time

- COMMITs
- replication
- "hint bits"
- statistics

IOPS + throughput == performance

optimizing IO

own HW: use SSDs, HW RAID

cloud: get more IOPS, EBS-optimized

also: No More EXT3! (or HFS or NTFS)

and: some Linux Kernel issues (3.2, 3.5)

RAM usage is thresholded

good: enough to cache "working set"


better: cache the whole database

best: DB fits in shared_buffers

some AWS tips

- use GP2 + extra GB
- PAAS != performance
- check zone distance

scaling infrastructure




the easy stuff

use the latest version of PostgreSQL
perf improvements in every release

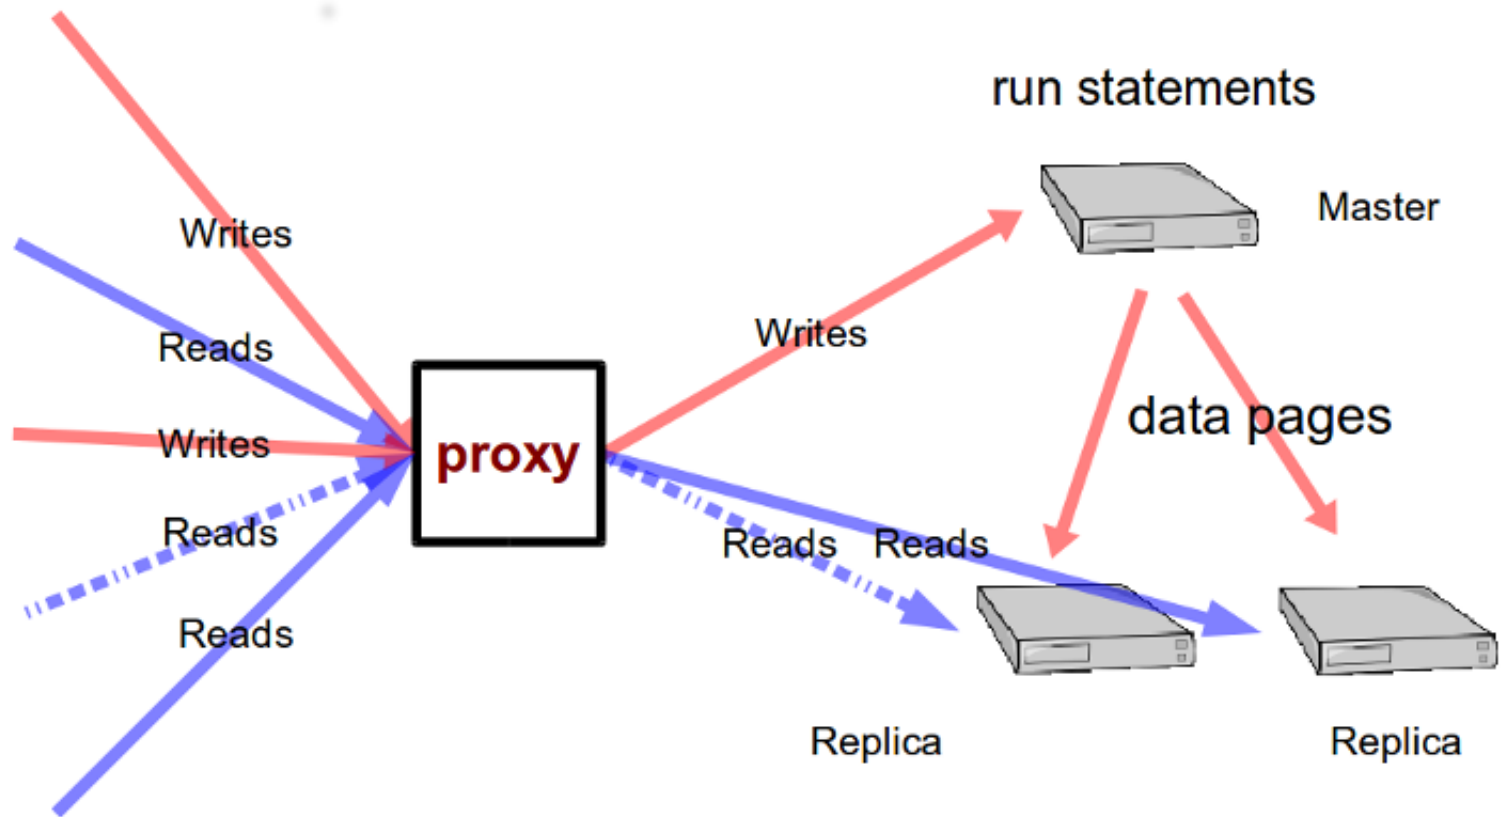
put PostgreSQL on its own
server/instance
databases do not share well

use pgbouncer for connection pooling

```
pool_mode = transaction
```



event-based pooling



load balance to read slaves

load balancing using routes

```
DATABASES = {  
    'master': {  
        'NAME': 'master',  
        'ENGINE': 'django.db.backends.pgsql'  
    },  
    'replica1': {  
        'NAME': 'replica1',  
        'ENGINE': 'django.db.backends.pgsql'  
    }  
    ...  
}
```

load balancing special workloads

- reporting
- cache refresh
- queueing (separate DB)

postgresql.conf

memory use

```
shared_buffers = 2GB
```

```
# RAM/4 up to 8GB
```

```
work_mem = 32MB
```

```
# 8MB to 32MB: web
```

```
# 128MB to 1GB: reporting
```

```
# limit: RAM/(max_connections/2)
```

more memory use

```
effective_cache_size = 6GB
```

```
# 3/4 of RAM
```

```
wal_buffers = 64MB
```

```
# just set it
```

```
maintenance_work_mem = 512MB
```

```
# RAM/32
```

```
# more for reporting
```

WAL

```
checkpoint_segments = 64  
# make WAL bigger  
# space / 32MB  
  
checkpoint_completion_target = 0.9
```

more settings

```
stats_temp_directory = '/mnt/ramdisk'
```

```
# helps with latency
```

```
random_page_cost = 1.5
```

```
# for AWS, SSD
```

```
effective_io_concurrency = 4
```

```
# for AWS, SSD, RAID
```


logging settings

```
log_connections = on  
log_disconnections = on  
log_temp_files = 1kB  
log_lock_waits = on  
log_checkpoints = on  
log_min_duration_statement = 0
```

recap

- Do Less querying
- fix resource-hungry requests
- get adequate hardware
- scale your infrastructure
- tune the config a little

questions?

more www.pgexperts.com
jberkus: www.databasesoup.com

more [austinPUG Thursday night:](http://austinPUG Thursday night)
events: www.meetup.com/austinpug

Postgres Open:
Sept 16-18, Dallas



pgConfSV: